

# The Agent Evaluation Framework

Where AI agents actually make sense (and where they don't)

Score any candidate process across five dimensions before writing a single line of agent code. Takes 10 minutes. Saves months.

- 1

**Trigger Clarity — Is it repeatable?**  
5 = fixed schedule / clear event   3 = mostly predictable   1 = ad-hoc  
Agents need clear triggers. If a human decides *when*, you haven't automated anything.
- 2

**Data Availability — Is the input structured?**  
5 = APIs / databases / structured files   3 = mixed   1 = tribal knowledge  
Gather then summarize. Pre-built scripts fetch data; the agent reads — it doesn't chase.
- 3

**Output Predictability — Is the decision space bounded?**  
5 = binary / small set of outcomes   3 = mostly bounded   1 = open-ended  
Agents excel at 'pick from options.' They fail at novel problem-solving.
- 4

**Risk Tolerance — What's the cost of error?**  
5 = low stakes / easy redo   3 = moderate, recoverable   1 = financial / legal risk  
Start where mistakes are cheap. Build trust. Expand scope.
- 5

**ROI Signal — What's the current time cost?**  
5 = 4+ hrs/week   3 = 1–3 hrs/week   1 = under 30 min/week  
If it takes 20 min/week, an agent isn't solving a problem — it's creating a maintenance obligation.

## Decision Matrix

| Total Score | Recommendation                                                                                           |
|-------------|----------------------------------------------------------------------------------------------------------|
| 21 – 25     | <b>Automate fully.</b> Textbook agent use case. Build it, deploy it, monitor it.                         |
| 15 – 20     | <b>Agent-assisted.</b> The agent drafts / categorizes — a human reviews. Most real-world wins live here. |
| 10 – 14     | <b>Not yet.</b> Restructure the process first (better inputs, clearer decisions) or the ROI isn't there. |
| Below 10    | <b>Keep human.</b> The agent would cause more problems than it solves. Revisit later.                    |

### Worked Example: Weekly Financial Reporting

A client's finance team spent 6 hrs/week compiling cash flow from three banks, two payment processors, and a spreadsheet.

| Dimension   | Score | Why                                       |
|-------------|-------|-------------------------------------------|
| Repeatable? | 5     | Every Monday, same trigger                |
| Structured? | 4     | APIs + one manual spreadsheet (fix first) |
| Bounded?    | 5     | Categorize, sum, flag anomalies           |
| Error cost? | 4     | Internal report, caught in review         |
| Time cost?  | 5     | 6 hrs/week of skilled time                |

**Total: 23 → Full automation.** Running in two weeks. The spreadsheet became a structured CSV. The agent pulls data Sunday night, generates the report, drops it in Slack Monday 7 AM. A human scans it in 10 min instead of building it in 6 hrs.

## How to Use This

1. List your top 5–10 repetitive processes   2. Score each one honestly   3. Sort by total   4. Start with the highest scorer   5. Ship in weeks, not months

*The bottleneck in AI adoption isn't the model. It's knowing where to point it.*