

The Agent FIT Scorecard

Is this process actually a good candidate for an AI agent?

Elliot Betancourt



Score any candidate process across six dimensions before you greenlight the project, not after.

- 1

Tool Surface

5 = one or two tools 3 = a handful, well-defined 1 = many bespoke integrations

Coding agents need four tools total: read, write, edit, run command.
- 2

Medium

5 = plain text / structured data 3 = mixed media 1 = physical world, voice, ambiguous docs

Text is cheap to generate, parse, and verify. Everything else costs more per step.
- 3

Correctness Signal

5 = automatic pass/fail check 3 = partial signal (linters, review) 1 = purely subjective

Coding agents self-correct because the compiler is a built-in critic.

- 4

Prior Art

5 = dense, widely documented 3 = some internal examples 1 = novel, undocumented territory

Decades of documented answers exist for parsing a CSV. Far fewer for your specific client.
- 5

Process Structure

5 = established methodology 3 = informal but consistent 1 = ad hoc every time

Software engineering already thinks in systems. Most operations work doesn't, yet.
- 6

Judgment Requirement

5 = none required 3 = some, but bounded 1 = high — judgment IS the value

Inverse-scored on purpose. High-judgment tasks are where humans stay the value, not the bottleneck.

Decision Matrix

Total Score	Recommendation
26-30	Strong candidate. Build the agent.
19-25	Partial candidate. Narrow the scope first.
12-18	Not yet. Restructure the process before automating it.
Below 12	Keep this one human.

The bottleneck usually isn't the model. It's picking the right candidate.