

Agent Performance Review

A one-page framework for evaluating AI agents

You've deployed agents. Now evaluate them. Rate each category 1–5, total out of 50. Run quarterly.

01 IDENTITY & ROLE

Name, function, domain, owner. If you can't explain it in one sentence, the purpose isn't clear.

02 RELIABILITY

Uptime, error rate, mean time to recovery. Measure both frequency and how fast it bounces back.

03 OUTPUT QUALITY

Accuracy, consistency, human override rate. Above 30% override = draft generator, not agent.

04 EFFICIENCY

Cost per task, tokens consumed, time to complete. Know the actual unit economics.

05 AUTONOMY LEVEL

Fully autonomous / supervised / manual trigger. Higher autonomy demands higher reliability.

06 RISK PROFILE

Failure impact, blast radius, recovery plan. Size the worst case honestly.

07 COLLABORATION

Handoff quality to humans and other agents. Can someone pick up where it left off?

08 GROWTH TRAJECTORY

Performance trend over 30/60/90 days. Agents that don't improve are replacement candidates.

09 TRUE COST OF OWNERSHIP

Infrastructure + API + human oversight hours. The API bill is the smallest line item.

10 RECOMMENDATION

Continue / Expand / Retrain / Replace / Retire. No agent is entitled to exist.

Total Score _____ / 50

Date _____

Reviewer _____

Agent _____

SCORING GUIDE

40–50

High Performer

30–39

Solid

20–29

Underperforming

Below 20

Replace / Retire