

The Agent Reliability Audit

Score your agent system against the three failure modes that don't crash

Elliot Betancourt

Most reliability work focuses on errors. The worst failures in agent systems don't throw any. They return success — and corrupt the downstream state until something catches it days later. This audit maps three detection layers to three failure modes. Score each dimension 5 (covered), 3 (partial), 1 (not in place).

1 Liveness — catches silent halts

5 = Heartbeat + alerting + recent verification 3 = Heartbeat exists, alerting incomplete 1 = "We'd notice if it stopped"

Silent halts: OAuth expires, rate limiter throttles to zero, the cron job misfires. Logs look healthy. Nothing's running. The fix is a deadman switch you actually watch — not a dashboard nobody opens.

2 Verification — catches polite lies

5 = Output validators + golden cases + pre-write guards 3 = Schema validators only 1 = "The LLM said it was done"

Polite lies: the API returns 200, the output parses, the agent reports success — and the semantics are subtly wrong. The fix is verifiers that test meaning, not shape. Plus pre-write guards on destructive operations through "helpful" APIs.

3 Drift Detection — catches drift

5 = Baselines + scheduled re-runs + drift alerts 3 = Baselines exist, no schedule 1 = "We'll know if it gets worse"

Drift: nothing in your code changed. The model behind it did. Same prompts, different answers. The fix is a fixed reference (golden cases) you re-run on a schedule, with alerts when output diverges past a threshold.

Decision Matrix

Total Score	Recommendation
13–15	Production-grade. Your agents fail loud.
9–12	Partial coverage. Audit your gaps.
5–8	Mostly trusting silence. High risk.
3–4	You're shipping silence. Stop and fix.

Loud failure is a gift. Quiet failure is a debt.