

The Agent Task Fit Matrix

5 Dimensions for Deciding Where AI Agents Belong

Elliot Betancourt

Most agent failures aren't model failures — they're deployment failures. Score any process across these five dimensions before you build.

1 Trigger Clarity — Does it know when to start?

5 = fixed schedule or unambiguous event 3 = mostly predictable, some edge cases 1 = ad-hoc, human decides when

If a human decides when to run it, you haven't automated anything — you've just added a step.

2 Output Verifiability — Can you check if it worked?

5 = binary outcome, testable without re-running 3 = format correct, quality needs review 1 = subjective — requires human judgment

Without verifiable output, you can't detect failure. The feedback loop stays open.

3 Error Cost — What happens when it's wrong?

5 = low cost, easily reversed 3 = moderate, correctable with effort 1 = high cost, hard to reverse

High-cost errors require human approval gates — which changes your entire architecture.

4 Tool Surface — How much can it touch?

5 = 1-2 tools 3 = 3-4 tools 1 = 5+ tools or broad API access

Each tool adds failure modes. Small, well-defined tool surfaces are dramatically more reliable.

5 Iteration Speed — How fast do you catch failures?

5 = immediate feedback 3 = same-day review cycle 1 = days to weeks

Slow feedback loops let failures compound. Fast loops let you improve.

Decision Matrix

Total Score	Recommendation
21–25	Strong fit. Automate fully.
15–20	Agent-assisted. Keep a human for edge cases.
10–14	Not ready. Restructure first, then revisit.
Below 10	Keep human. The process needs work.

The bottleneck isn't the model. It's knowing where to point it.