

The AI Agent Systems Audit

5 Dimensions to Diagnose Why Your AI Project Isn't Working

Elliot Betancourt

The first instinct when an AI project stalls is to blame the model. Nine times out of ten, the model isn't the problem. Score each dimension 5, 3, or 1. The total (out of 25) tells you where to invest next.

1 Architecture Clarity — Can you draw your system on a napkin?

5 = documented, named components, explicit interfaces 3 = informal design in someone's head 1 = structure emerged accidentally around a model choice

Without an explicit architecture, you cannot reason about tradeoffs. You are shipping whatever the LLM suggests.

2 Model Right-Sizing — Paying Opus for Haiku work?

5 = model tier matched to task; cost per task measured 3 = one model for everything; cost is 'fine' 1 = defaulted to the biggest model; no measurement

The biggest model is almost never the right model. Cost scales with intent, not defaults.

3 Determinism Layer — Is the LLM doing shell-script work?

5 = retries, state, routing, fetching are all deterministic code 3 = LLM does some orchestration it shouldn't 1 = agent is mostly an LLM with vibes; no scaffold

LLMs reason. Scripts fetch. Code handles the rest. Most projects fail at this dimension.

4 Agent Specialization — One mega-agent or a team?

5 = multiple specialized agents with bounded scope 3 = one general agent with multiple tool categories 1 = one mega-agent trying to do everything

A team of specialists beats a general-purpose agent. Each knows its job, none knows everything.

5 Observability — Can you explain a decision from yesterday?

5 = decisions logged with reasoning; memory persists across runs 3 = logs exist but no decision trace or cross-run memory 1 = no logs; cannot reproduce yesterday's behavior

If you cannot see what the agent decided and why, you cannot improve it. Memory is architectural, not a model capability.

Decision Matrix

Total Score	Recommendation
21–25	Ship it. Your system is mature — invest in scale.
15–20	Tighten your lowest score before adding features.
10–14	Stop shipping. Redesign weak dimensions first.
Below 10	Back to architecture. The system isn't architected yet.

Most AI projects fail at dimension 3. The LLM isn't the bottleneck — the scaffold around it is. The model is commodity. Architecture is the moat.