

AI Systems Readiness Checklist

6 Questions That Separate the 12% That Ship from the 88% That Don't

Elliot Betancourt

Most AI agent projects fail not because the model isn't good enough, but because the system around it was never designed to stay coherent over time. Score each dimension to find out if your project is model-centric or systems-ready.

1 Error Recovery

- Defined fallbacks and retry logic
- Manual intervention, but documented
- Undefined — workflow just stops

Why it matters: A 10-step chain with no recovery is a 10x failure multiplier.

2 State Management

- Clean boundaries between steps
- Shared state, occasional conflicts
- Global state — everything sees everything

Why it matters: State leakage is the #1 cause of agents degrading over time.

3 Human Handoff

- Clear escalation triggers defined
- We'll figure it out when it happens
- No handoff plan at all

Why it matters: Agents without exit ramps run until they fail visibly.

4 Accuracy Compounding

- End-to-end accuracy measured
- Per-step metrics only
- The model is 95% accurate (stopped there)

Why it matters: 85% per step across 10 steps = 20% end-to-end.

5 Memory Architecture

- Tiered context: raw, structured, summary
- Conversation history as only memory
- No memory beyond context window

Why it matters: Long-running agents need managed memory, not just more tokens.

6 Cost at Scale

- Measured at 10x volume with tiering
- Dev-scale numbers only
- Haven't measured

Why it matters: Demo costs are fiction. Production costs are the business case.

SCORING GUIDE



6 Green:

Ship it. Your architecture supports the model.



3+ Yellow:

Fixable. Focus yellows before scaling.



Any Red:

Stop. Fix the system before upgrading the model.

The 88% that fail aren't using bad models. They're running good models through bad systems.