

# AI Unit Economics Checklist

## 6 Questions to Ask Before Greenlighting Any AI Feature

Elliot Betancourt

OpenAI reportedly burned \$15M/day running Sora before anyone asked whether the math worked — downloads dropped 66% in three months and Disney walked away without spending a dollar. The tech wasn't the problem; the economics were. This checklist is what I run on every AI project I touch: six questions, scored on a traffic-light system, designed to surface cost risks before they become existential.

### 1 What does a single inference call cost at the volume your users will actually generate?

- Measured cost per request at 10x current volume
- Dev/test numbers only — haven't modeled production scale
- Quoting demo costs or haven't measured at all

*Why it matters:* Demo-scale costs are fiction. \$0.02/request at 100 users can become \$0.15 at 10,000.

### 2 Does each new user add inference cost, or does your architecture amortize it?

- Marginal cost is flat or declining (caching, batching, pre-computation)
- Linear cost growth — each user adds proportional inference cost
- Super-linear — costs accelerate with scale (per-user fine-tuning, real-time gen)

*Why it matters:* Linear is survivable. Super-linear is a death sentence at scale.

### 3 Does the value delivered to the user justify what you spend on each call?

- Clear ROI — feature saves time/money worth 5x+ the inference cost
- Value is real but hard to quantify, or margin is thin (2–3x)
- Feature is "cool" but doesn't map to measurable user value

*Why it matters:* This is where Sora died. Stunning output, but no one could justify the cost per generation.

### 4 Are you using the smallest model that delivers acceptable quality?

- Tested multiple model sizes, chose the smallest that meets quality bar
- Using whatever model you prototyped with — haven't tested alternatives
- Using the largest available model "to be safe"

*Why it matters:* Cost difference between largest and smallest models from the same provider is often 60x.

### 5 What's your plan if your AI provider doubles pricing overnight?

- Fallback provider tested in staging; abstraction layer supports switching
- Know alternatives exist but haven't tested them
- Locked to one provider with no switching plan

*Why it matters:* AI pricing is not stable. A 2x price hike shouldn't be an existential crisis.

### 6 Does your revenue model absorb inference costs, or are you subsidizing every interaction?

- Inference cost is a known line item in pricing (usage-based, tiered, or built into margin)
- Cost exists but is buried in flat subscription — margin depends on usage patterns
- Eating inference cost to drive adoption, with no path to cost recovery

*Why it matters:* "We'll figure out monetization later" doesn't work for costs that scale with usage.

## SCORING GUIDE

#### ● 6 Green

Ship it — you've done the homework.

#### ● 4–5 Green

Proceed with caution. Address yellows before scaling.

#### ● Any Red

Stop and fix before committing more resources.

#### ● 3+ Red

Kill the project or fundamentally rethink.