

Agentic AI Governance Readiness Checklist

6 Questions to Ask Before Greenlighting Any Agent Project

Elliot Betancourt

Gartner predicts over 40% of agentic AI projects will be canceled by 2027 — not because the models fail, but because of escalating costs, unclear ROI, and missing risk controls. This is the six-item check I run before I'll say yes to a client's agentic AI project.

1 Decision Authority

- Named owner with override authority
- Informal, no single owner
- No one has been asked

Why it matters: No owner means no one stops it in time.

2 Audit Trail

- Full logging to a reviewable record
- Partial logging
- No logging beyond app defaults

Why it matters: You can't fix what you can't trace.

3 Cost Ceiling

- Hard limit enforced in code
- Soft budget, monitored manually
- No ceiling defined

Why it matters: Escalating cost is Gartner's #1 cancellation reason.

4 Human Escalation Path

- Specific conditions documented and tested
- "Someone will notice"
- No path defined

Why it matters: Hope is not an escalation policy.

5 Data Access Scope

- Scoped to task, reviewed
- Broader than needed but bounded
- Full account or API access

Why it matters: Blast radius is a design decision, not an accident.

6 Rollback Plan

- Tested kill switch
- Theoretically possible
- Never tested

Why it matters: An untested rollback is not a rollback.

SCORING GUIDE



6 Green:

Ship it.



1-2 Yellow:

Fix before launch, not after.



Any Red:

Stop. This is the failure mode Gartner is pricing at 40%.

The best time to run this is before you write the first line of code.